

# **Homework-3**

CSCE 633: Machine Learning

Manav Gurumoorthy

830000011

## Introduction:

In this HW, we looked at the use of Artificial Neural Networks to perform image recognition.

The dataset consists of labelled images of humans depicting different facial expression. The recognition task was to accurately discern the facial expression being depicted in the image.

The dataset was consisted of several 48X48 grayscale images. i.e. each image was considered as a vector of length 2304, with each element representing the intensity of that pixel.

We look at two main approaches here

1. The use of a Feed Forward Network
2. Use of a Convolutional Neural Network.

While these are the two broad approaches in this study, each method had their own nuances and required careful manipulation of hyperparameters.

## Data Exploration:

In this part, we explored the data that we were going to model. The first step was to plot the images from the data that we had. This was done using inbuilt functions in python.



(one such example) more in the code section.

We also looked at how many exhibits from each of the labels were present in the training dataset.

Angry:	3995
Disgust:	436
Fear:	4097
Happy:	7215
Sad:	4830
Surprise:	3171
Neutral:	4965

---

Total:	28709
--------	-------

## Feedforward Neural Network:


In this section we construct a FNN using the keras library and test the performance of the model with the validation data by playing around with hyperparameter values. I liked the hyperopt library, so I used this to find optimal hyperparameter values. My search space consisted of the following parameters:

```

space = {
    # The activation_fn choices:
    'activation_fn': hp.choice('activation_fn', ['relu']),
    # Uniform distribution in finding appropriate dropout values
    'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.5),
    # Choice of optimizer
    'optimizer': hp.choice('optimizer', ['Adam', 'sgd']),
}

```

We also plotted the loss as a function of the number of iterations (epochs?). One such example is shown below. The rest are shown as a part of the code section.



Some stats that are printed with every change in hyperparameters: (This was done using Prafulla's code)

```

Total Training Time is (s):
7.4364213943481445
Hyperparameters:
{'activation_fn': 'relu', 'dropout_prob': 0.25336800158266903, 'optimizer': 'sgd'}
Accuracy:
0.39091668988991396
-----
4% | 1/25 [00:07<03:11, 8.00s/it, best loss: -0.39091668988991396]
Model loss

```

The best hyperparams were then used and the model was tested on the test set.

```

=====
Best Hyperparameters {'activation_fn': 0, 'dropout_prob': 0.2476540471215125, 'optimizer': 0}
3589/3589 [=====] - 0s 49us/step
=====
Test Accuracy: 0.45528002229863535

```

In an earlier test I also tried 3 different activation functions namely ReLu, Sigmoid and tanh, but since the other 2 gave me extremely poor results I defaulted to just using ReLu.

## Convolution Neural Networks:

Like the FNN case, we create a CNN model using the keras library APIs. Similar to the FNN case, we try to optimise the model using the validation set and different combinations of hyperparameters to optimise the model.

Firstly I used a simple for loop which used the following params and trained 10 different models, the best hyperparameter pair from this was used on the test set.

```

1 dropout_range = [0.1,0.2,0.3,0.4,0.5]
2 kernels = [2,3]

```

```

Best Feature Pair is: 3 and 0.3
Best performance from training: 0.45082195598489894
Kernel Size: 3
Dropout: 0.3

```

```
3589/3589 [=====] - 2s 522us/step
Accuracy on Test samples: 0.2449150181191982
```

Unfortunately, this was not yielding very good results because I was only using a small set of values. So once again I used HyperOpt to find a better set of hyperparameters.

```
-----
100%|██████████| 25/25 [12:46<00:00, 30.65s/it, best loss: -0.5937587071939842]
=====
Best Hyperparameters {'conv_kernel_size': 1, 'dropout_prob': 0.2890305052408365, 'optimizer': 0}
3589/3589 [=====] - 0s 102us/step
=====
Test Accuracy: 0.5851212036862085
```

## Fine Tuning:

I was a little unclear about this part so I have done 2 things.

The first thing is to train a model on the MNIST dataset and modify the FER dataset suitably to work with the trained model. Doing this got an accuracy of 40%

```
Model: "sequential_29"
Layer (type)                 Output Shape                 Param #
-----
conv2d_9 (Conv2D)            (None, 26, 26, 32)         320
conv2d_10 (Conv2D)           (None, 24, 24, 32)         9248
max_pooling2d_5 (MaxPooling2 (None, 12, 12, 32)         0
dropout_31 (Dropout)         (None, 12, 12, 32)         0
conv2d_11 (Conv2D)           (None, 10, 10, 64)         18496
conv2d_12 (Conv2D)           (None, 8, 8, 64)           36928
max_pooling2d_6 (MaxPooling2 (None, 4, 4, 64)           0
dropout_32 (Dropout)         (None, 4, 4, 64)           0
flatten_3 (Flatten)          (None, 1024)                0
dense_31 (Dense)              (None, 512)                 524800
dense_32 (Dense)              (None, 10)                  5130
-----
Total params: 594,922
Trainable params: 594,922
Non-trainable params: 0
```

Then adding another layer for the outputs of FER,

```
Model: "sequential_29"
Layer (type)                 Output Shape                 Param #
-----
conv2d_9 (Conv2D)            (None, 26, 26, 32)         320
conv2d_10 (Conv2D)           (None, 24, 24, 32)         9248
max_pooling2d_5 (MaxPooling2 (None, 12, 12, 32)         0
dropout_31 (Dropout)         (None, 12, 12, 32)         0
conv2d_11 (Conv2D)           (None, 10, 10, 64)         18496
conv2d_12 (Conv2D)           (None, 8, 8, 64)           36928
max_pooling2d_6 (MaxPooling2 (None, 4, 4, 64)           0
dropout_32 (Dropout)         (None, 4, 4, 64)           0
flatten_3 (Flatten)          (None, 1024)                0
dense_31 (Dense)              (None, 512)                 524800
dense_32 (Dense)              (None, 10)                  5130
dense_35 (Dense)              (None, 7)                   77
-----
Total params: 659,914
Trainable params: 594,922
Non-trainable params: 64,992
```

The final model accuracy:

```
3589/3589 [=====] - 1s 232us/step
Accuracy on Test samples: 0.4093062134299248
```

The second thing was to create a combined model of FER and MNIST.

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 28, 28, 1)	0	
input_2 (InputLayer)	(None, 48, 48, 1)	0	
sequential_2 (Sequential)	multiple	98272	input_1[0][0] input_2[0][0]
MNIST (Dense)	(None, 10)	81930	sequential_2[1][0]
FER (Dense)	(None, 7)	290311	sequential_2[2][0]
Total params: 470,513			
Trainable params: 470,513			
Non-trainable params: 0			

This yielded the following results,

```
===
MNIST Accuracy: 0.9875731395127282
FER Accuracy: 0.8625562360739103
```

## Data Augmentation:

In this section we look at different types of data augmentation using the ImageDataGenerator from keras. I tried Standardisation, ZCA Whitening and Random Rotation, the results are shown below in the same order.

```
=====
Best Hyperparameters {'conv_kernel_size': 2, 'dropout_prob': 0.21040384851779753, 'optimizer': 0}
3589/3589 [=====] - 0s 92us/step
=====
Test Accuracy: 0.6032320980816108
```

```
=====
Best Hyperparameters {'conv_kernel_size': 2, 'dropout_prob': 0.14865966844802445, 'optimizer': 0}
3589/3589 [=====] - 0s 112us/step
=====
Test Accuracy: 0.5918083031568131
```

```
=====
Best Hyperparameters {'conv_kernel_size': 2, 'dropout_prob': 0.19430122651964926, 'optimizer': 0}
3589/3589 [=====] - 0s 123us/step
=====
Test Accuracy: 0.5770409584925612
```

## Feature Design:

In this section I tried implementing HOG on the images because it seemed like it would give good results for the task we had intended. Unfortunately, I guess either there is an error in my implementation or HOG is not a method that is suitable for this sort of task. But I attribute it to the former rather than the latter.

```
-----
Total Training Time is (s):
26.25912308692932
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.10813003008178762, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
21.570691347122192
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.27555616085888457, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
100%|██████████| 25/25 [10:50<00:00, 26.00s/it, best loss: -0.24937308442878273]
=====
Best Hyperparameters {'conv_kernel_size': 2, 'dropout_prob': 0.15049653462509077, 'optimizer': 1}
3589/3589 [=====] - 0s 87us/step
=====
Test Accuracy: 0.2449150181191982
```

## APPENDIX A: COLAB Workbook

## ▼ DATA EXPLORATION

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

☞ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.m

```
1 #IMPORTING THE DATA FROM THE DATASETS.
2 import pandas as pd
3 import numpy as np
4 import csv
5 import matplotlib.pyplot as plt
6 from skimage import data, io, filters
7
8 path = "/content/drive/My Drive/Homework3_Data/Train_Data.csv"
9
10 path3 = "/content/drive/My Drive/Homework3_Data/Validation_Data.csv"
11 train_data_df = pd.read_csv(path)
12
13 validation_data_df = pd.read_csv(path3)
```

```
1 #Creating data arrays
2 train_labels = train_data_df['emotion']
3 valid_labels = validation_data_df['emotion']
4
5 train_pixels = train_data_df.drop('emotion',axis=1)
6 train_pixels = train_pixels.to_numpy()
7
8 valid_pixels = validation_data_df.drop('emotion',axis=1)
9 valid_pixels = valid_pixels.to_numpy()
10
11 def image_display(i):
12     image = train_pixels[i][0]
13     shape = (48,48)
14     image = [int(k) for k in image.split(' ')]
15     image = np.array(image)
16     image = image.reshape(shape)
17     plt.title(train_labels[i])
18     plt.imshow(image)
19     plt.show()
```

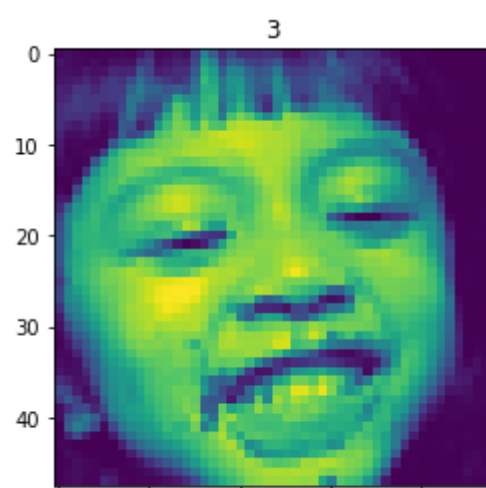
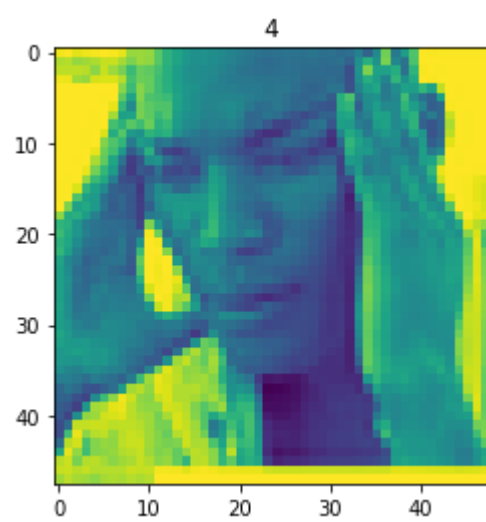


```
1 #-----TESTTING STUFF
2 path2 = "/content/drive/My Drive/Homework3_Data/Test_Data.csv"
3 test_data_df = pd.read_csv(path2)
4 test_labels = test_data_df['emotion']
5
6 test_pixels = test_data_df.drop('emotion',axis=1)
7 test_pixels = test_pixels.to_numpy()
8 test_image_array = []
9 for i in range (len(test_pixels)):
```


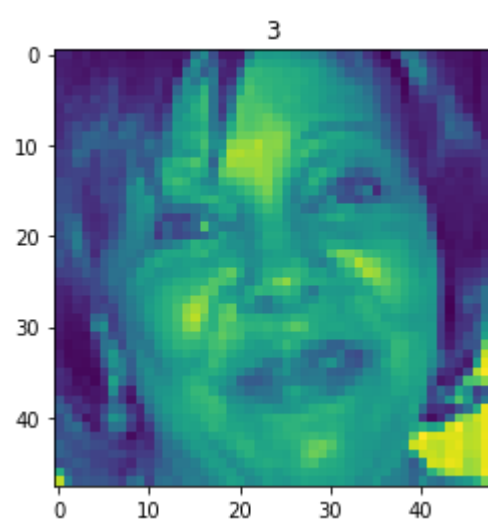





```
10 temp = test_pixels[i][0]
11 temp = [int(k) for k in temp.split(' ')]
12 temp = np.array(temp)
13 test_image_array.append(temp)
14
15 test_image_array = np.array(test_image_array)
16
```

```
1 import random
2 my_randoms=[]
3 for i in range (10):
4     my_randoms.append(random.randrange(1,len(train_pixels),1))
5 for i in my_randoms:
6     image_display(i)
```

☞







```

1 count0 = 0
2 count1 = 0
3 count2 = 0
4 count3 = 0
5 count4 = 0
6 count5 = 0
7 count6 = 0
8 for i in range(len(train_labels)):
9     if (train_labels[i]==0):
10         count0 += 1
11     elif(train_labels[i]==1):
12         count1 += 1
13     elif(train_labels[i]==2):
14         count2 += 1
15     elif(train_labels[i]==3):
16         count3 += 1
17     elif(train_labels[i]==4):
18         count4 += 1
19     elif(train_labels[i]==5):
20         count5 += 1
21     elif(train_labels[i]==6):
22         count6 += 1
23     else:
24         print("Done ")

```

```

24     print( 'oops. ' )
25 total = count0+count1+count2+count3+count4+\
26         count5+count6
27 print("\nAngry:   ",count0,"\nDisgust:  ", count1,"\nFear:   ",count2,\
28       "\nHappy:   ",count3,"\nSad:     ",count4,"\nSurprise:" \
29       , count5,"\nNeutral: " ,count6,"\n_____\nTotal:  ",total)

```



```

Angry:    3995
Disgust:   436
Fear:     4097
Happy:    7215
Sad:      4830
Surprise: 3171
Neutral:  4965

Total:    28709

```

## ▼ FNN

```

1 train_image_array = []
2 for i in range (len(train_pixels)):
3     temp = train_pixels[i][0]
4     temp = [int(k) for k in temp.split(' ')]
5     temp = np.array(temp)
6     train_image_array.append(temp)

```

```

1 valid_image_array = []
2 for i in range (len(valid_pixels)):
3     temp = valid_pixels[i][0]
4     temp = [int(k) for k in temp.split(' ')]
5     temp = np.array(temp)
6     valid_image_array.append(temp)

```

```

1 train_image_array = np.array(train_image_array)
2 valid_image_array = np.array(valid_image_array)

```

```

1 from skimage import data, io, filters
2 #normalize all the images
3 train_image_array = (train_image_array / 255) - 0.5
4 valid_image_array = (valid_image_array / 255) -0.5
5 test_image_array = (test_image_array / 255) -0.5
6
7 # train_image_array = filters.sobel(train_image_array)
8 # valid_image_array = filters.sobel(valid_image_array)
9 # test_image_array = filters.sobel(test_image_array)

```

```

1 import warnings
2 import keras
3 warnings.filterwarnings("ignore")
4 from keras.models import Sequential
5 from keras.layers import Dense, Dropout

```

```

6 from keras.utils import to_categorical
7 from keras import regularizers
8
9
10 # # Flatten the images into vectors (1D) for feed forward network
11 # flatten_train_images = train_image_array.reshape((-1, 48*48))
12 # flatten_test_images = test_image_array.reshape((-1, 48*48))
13 # flatten_valid_images = valid_image_array.reshape((-1, 48*48))

```

☞ Using TensorFlow backend.

The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.

We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the %tensorflow\_version --tf-gpu magic.

```

1 import time
2 class TimeHistory(keras.callbacks.Callback):
3     def on_train_begin(self, logs={}):
4         self.times = []
5
6     def on_epoch_begin(self, batch, logs={}):
7         self.epoch_time_start = time.time()
8
9     def on_epoch_end(self, batch, logs={}):
10         self.times.append(time.time() - self.epoch_time_start)

```

```

1 # Evaluate your model's performance on the test data
2 def model_test(model):
3     performance = model.evaluate(valid_image_array, to_categorical(valid_labels))
4     print("\nAccuracy on Test samples: {0}".format(performance[1]))

```

```

1 # Compiling the model
2 def model_compile(model):
3     model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'],
4 # Train model
5     time_callback = TimeHistory()
6     hist = model.fit(train_image_array, to_categorical(train_labels), epochs=10, batch_size=128,
7
8     plt.plot(hist.history['loss'])
9     plt.title('Model loss')
10    plt.ylabel('Loss')
11    plt.xlabel('Epoch')
12    plt.legend(['Train'], loc='upper left')
13    plt.show()
14    print("\nTotal Training Time is: ", sum(time_callback.times), 's.\n')

```

```

1 from hyperopt import hp, fmin, tpe, STATUS_OK, Trials
2
3 #MODEL 1
4 #Define a Feed-Forward Model with 2 hidden layers with dimensions 392 and 196 Neurons
5 def optimize_fnn(hyperparameter):
6     model = Sequential([
7         Dense(2304, activation=hyperparameter['activation_fn'], input_shape=(48*48,), name='first_hidden_layer'),
8         Dense(2304//2, activation=hyperparameter['activation_fn'], name="second_hidden_layer"),
9         Dense(7, activation='softmax')

```

```


9     sense(, activation = softmax ,
10 ])
11 # Validate your Model Architecture
12 #print(model1.summary())
13 model.compile(optimizer=hyperparameter['optimizer'], loss='categorical_crossentropy',
14               time_callback = TimeHistory())
15 hist = model.fit(train_image_array, to_categorical(train_labels), epochs=10, batch_si
16 plt.plot(hist.history['loss'])
17 plt.title('Model loss')
18 plt.ylabel('Loss')
19 plt.xlabel('Epoch')
20 plt.legend(['Train'], loc='upper left')
21 plt.show()
22 print("Total Training Time is (s): ", sum(time_callback.times))
23
24 performance = model.evaluate(valid_image_array, to_categorical(valid_labels), verbose
25
26 print("Hyperparameters: ", hyperparameter, "Accuracy: ", performance[1])
27 print("-----")
28 # We want to minimize loss i.e. negative of accuracy
29 return({"status": STATUS_OK, "loss": -1*performance[1], "model":model})
30
31 # Define search space for hyper-parameters
32 space = {
33     # The activation_fn choices:
34     'activation_fn':hp.choice('activation_fn',['relu']),
35     # Uniform distribution in finding appropriate dropout values
36     'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.5),
37     # Choice of optimizer
38     'optimizer': hp.choice('optimizer', ['Adam', 'sgd']),
39 }
40
41 trials = Trials()
42
43 best = fmin(
44     optimize_fnn,
45     space,
46     algo=tpe.suggest,
47     trials=trials,
48     max_evals=25,
49 )
50
51 print("=====")
52 print("Best Hyperparameters", best)
53
54 # You can retrain the final model with optimal hyperparameters on train+validation data
55
56 # Or you can use the model returned directly
57 # Find trial which has minimum loss value and use that model to perform evaluation on t
58 test_model = trials.results[np.argmin([r['loss'] for r in trials.results])]['model']
59
60 performance = test_model.evaluate(test_image_array, to_categorical(test_labels))
61
62 print("=====")
63 print("Test Accuracy: ", performance[1])
64

```





0%| | 0/25 [00:00<?, ?it/s, best loss: ?]



Total Training Time is (s):

7.4364213943481445


Hyperparameters:

`{'activation_fn': 'relu', 'dropout_prob': 0.25336800158266903, 'optimizer': 'sgd'}`

Accuracy:

0.39091668988991396

4%| | 1/25 [00:07<03:11, 8.00s/it, best loss: -0.39091668988991396]



Total Training Time is (s):

9.509779930114746


Hyperparameters:


`{'activation_fn': 'relu', 'dropout_prob': 0.10552055620174437, 'optimizer': 'Adam'}`

Accuracy:

0.4488715519726393

8%| | 2/25 [00:18<03:19, 8.66s/it, best loss: -0.4488715519726393]





Total Training Time is (s):

7.328493356704712


Hyperparameters:

`{'activation_fn': 'relu', 'dropout_prob': 0.1759849523388486, 'optimizer': 'sgd'}`

Accuracy:

0.38311507384502713

12% | 3/25 [00:26<03:05, 8.44s/it, best loss: -0.4488715519726393]



Total Training Time is (s):

9.538674116134644


Hyperparameters:

`{'activation_fn': 'relu', 'dropout_prob': 0.18842808735315886, 'optimizer': 'Adam'}`

Accuracy:

0.4555586514681553

16% | 4/25 [00:36<03:08, 8.98s/it, best loss: -0.4555586514681553]



Total Training Time is (s):

9.489769220352173


Hyperparameters:

`{'activation_fn': 'relu', 'dropout_prob': 0.2476540471215125, 'optimizer': 'Adam'}`

Accuracy:

0.4636388966368912

20% | 5/25 [00:46<03:07, 9.35s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

9.495479583740234


Hyperparameters:

`{'activation_fn': 'relu', 'dropout_prob': 0.3644742571285925, 'optimizer': 'Adam'}`

Accuracy:

0.44859292282803076

24% |██████████| 6/25 [00:56<03:02, 9.62s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

7.4017863273620605


Hyperparameters:


`{'activation_fn': 'relu', 'dropout_prob': 0.2555559787727395, 'optimizer': 'sgd'}`

Accuracy:

0.3842295904234612

28% |██████████| 7/25 [01:04<02:44, 9.15s/it, best loss: -0.4636388966368912]





Total Training Time is (s):

9.604174613952637


Hyperparameters:

`{'activation_fn': 'relu', 'dropout_prob': 0.32319345246869036, 'optimizer': 'Adam'}`

Accuracy:

0.45277236002207

-----  
32%|██████| | 8/25 [01:15<02:42, 9.57s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

7.4649622440338135


Hyperparameters:

`{'activation_fn': 'relu', 'dropout_prob': 0.3406336734941001, 'optimizer': 'sgd'}`

Accuracy:

0.38255781556411383

-----  
36%|██████| | 9/25 [01:23<02:26, 9.15s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

9.565871000289917

Hyperparameters:


`{'activation_fn': 'relu', 'dropout_prob': 0.4370333627480284, 'optimizer': 'Adam'}`

Accuracy:

0.44998606855107337

-----  
40%|██████| | 10/25 [01:34<02:22, 9.52s/it, best loss: -0.4636388966368912]

Model loss



Total Training Time is (s):

7.352118968963623


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.2464429141873418, 'optimizer': 'sgd'}

Accuracy:

0.3836723321342442

44%|██████| 11/25 [01:42<02:07, 9.10s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

9.775615215301514


Hyperparameters:


{'activation\_fn': 'relu', 'dropout\_prob': 0.1675727368174647, 'optimizer': 'Adam'}

Accuracy:

0.4533296182863757

48%|██████| 12/25 [01:52<02:04, 9.56s/it, best loss: -0.4636388966368912]





Total Training Time is (s):

9.897839069366455


Hyperparameters:

`{'activation_fn': 'relu', 'dropout_prob': 0.3479228659873571, 'optimizer': 'Adam'}`

Accuracy:

0.4488715519726393

-----  
52%|██████| | 13/25 [02:03<01:59, 9.93s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

7.853139162063599


Hyperparameters:

`{'activation_fn': 'relu', 'dropout_prob': 0.23118676879866445, 'optimizer': 'sgd'}`

Accuracy:

0.3898021733156317

-----  
56%|██████| | 14/25 [02:12<01:45, 9.55s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

7.92191481590271


Hyperparameters:


`{'activation_fn': 'relu', 'dropout_prob': 0.11349796190538407, 'optimizer': 'sgd'}`

Accuracy:

0.3870158818695465

-----  
60%|██████| | 15/25 [02:20<01:33, 9.31s/it, best loss: -0.4636388966368912]





Total Training Time is (s):

10.156506061553955


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.10106341569977983, 'optimizer': 'Adam'}

Accuracy:

0.44051267765929486

64% |██████████ | 16/25 [02:32<01:28, 9.83s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

8.006651639938354


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.49276434949032477, 'optimizer': 'sgd'}

Accuracy:

0.38478684870852636

68% |██████████ | 17/25 [02:40<01:16, 9.55s/it, best loss: -0.4636388966368912]



Epoch

Total Training Time is (s):

8.145221471786499


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.11279906259430761, 'optimizer': 'sgd'}

Accuracy:

0.3825578155516582

72% |██████████| | 18/25 [02:50<01:06, 9.44s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

10.069785594940186


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.21733738187782758, 'optimizer': 'Adam'}

Accuracy:

0.45834494288517724

76% |██████████| | 19/25 [03:01<00:59, 9.93s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

10.192638874053955


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.14335957639670258, 'optimizer': 'Adam'}


Accuracy:

0.4589022011785462

80% |██████████| | 20/25 [03:12<00:51, 10.32s/it, best loss: -0.4636388966368912]







Total Training Time is (s):

10.345960855484009


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.2936624860958402, 'optimizer': 'Adam'}

Accuracy:

0.45110058512950746

84% |██████████ | 21/25 [03:23<00:42, 10.64s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

10.243130207061768


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.15123480999583044, 'optimizer': 'Adam'}

Accuracy:

0.4494288102618563

88% |██████████ | 22/25 [03:35<00:32, 10.84s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

10.352617740631104


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.28900076526664475, 'optimizer': 'Adam'}

Accuracy:

0.44524937309272844

-----  
92%|██████████ | 23/25 [03:46<00:22, 11.03s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

10.280369758605957


Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.40388349489252084, 'optimizer': 'Adam'}

Accuracy:

0.45082195600981034

-----  
96%|██████████ | 24/25 [03:57<00:11, 11.13s/it, best loss: -0.4636388966368912]



Total Training Time is (s):

10.576332807540894

Hyperparameters:

{'activation\_fn': 'relu', 'dropout\_prob': 0.2051775684138893, 'optimizer': 'Adam'}

Accuracy:

0.45277236002207

-----  
100%|██████████ | 25/25 [04:09<00:00, 9.98s/it, best loss: -0.4636388966368912]

=====

Best Hyperparameters {'activation\_fn': 0, 'dropout\_prob': 0.2476540471215125, 'optimi  
3589/3589 [=====] - 0s 49us/step

=====

Test Accuracy: 0.45528002229863535

## ▼ CNNs

```
1 from keras.layers import Conv2D, Flatten, MaxPooling2D
2
3 # Flatten the images into vectors (1D) for feed forward network
4 train_images_3d = train_image_array.reshape((len(train_image_array),48,48,1))
5 test_images_3d = test_image_array.reshape((len(test_image_array), 48,48,1))
6 valid_images_3d = valid_image_array.reshape((len(valid_image_array), 48,48,1))
7
8 # Define 2 groups of layers: features layer (convolutions) and classification layer
9 common_features = [Conv2D(64, kernel_size=(3,3), strides=3, activation='relu', input_shape=(48,48,1)),
10                   Conv2D(64, kernel_size=(3,3), activation='relu'),
11                   MaxPooling2D(pool_size=(2,2)), Dropout(0.3),
12                   Conv2D(128, kernel_size=3, activation='relu'),
13                   Conv2D(128, kernel_size=3, activation='relu'),
14                   MaxPooling2D(pool_size=(2,2)), Dropout(0.3), Flatten(),]
15
16 classifier = [Dense(512, activation='relu'), Dense(7, activation='relu'),]
17 cnn_model = Sequential(common_features+classifier)
18 print(cnn_model.summary())
19 cnn_model.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['accuracy'])
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorf1

Model: "sequential\_27"

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 16, 16, 64)	640
conv2d_2 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_1 (MaxPooling2	(None, 7, 7, 64)	0
dropout_27 (Dropout)	(None, 7, 7, 64)	0
conv2d_3 (Conv2D)	(None, 5, 5, 128)	73856
conv2d_4 (Conv2D)	(None, 3, 3, 128)	147584
max_pooling2d_2 (MaxPooling2	(None, 1, 1, 128)	0
dropout_28 (Dropout)	(None, 1, 1, 128)	0
flatten_1 (Flatten)	(None, 128)	0
dense_27 (Dense)	(None, 512)	66048
dense_28 (Dense)	(None, 7)	3591
=====		
Total params: 328,647		
Trainable params: 328,647		
Non-trainable params: 0		
None		

```
1 from keras.layers import Conv2D, Flatten, MaxPooling2D
2 def run_cnn_model(kSize, dropout_rate):
3     # Define 2 groups of layers: features layer (convolutions) and classification layer
4     common_features = [Conv2D(64, kernel_size=(kSize), strides=3, activation='relu', input_shape=(3, 32, 32)),
5                         Conv2D(64, kernel_size=(kSize), activation='relu'),
6                         MaxPooling2D(pool_size=(2,2)), Dropout(dropout_rate),
7                         Conv2D(128, kernel_size=(kSize), activation='relu'),
8                         Conv2D(128, kernel_size=(kSize), activation='relu'),
9                         MaxPooling2D(pool_size=(2,2)), Dropout(dropout_rate), Flatten(),]
10
11     classifier = [Dense(512, activation='relu'), Dense(7, activation='relu'),]
12     cnn_model = Sequential(common_features+classifier)
13     print("Kernel Size: {0} \nDropout: {1}".format(kSize,dropout_rate))
14     print(cnn_model.summary())
15     cnn_model.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['accuracy'])
16     # Train model
17     cnn_model.fit(train_images_3d, to_categorical(train_labels), epochs=10, batch_size=256)
18
19     performance = cnn_model.evaluate(valid_images_3d, to_categorical(valid_labels))
20     print("Accuracy on Test samples: {0}".format(performance[1]))
21     results.append(performance[1])
```

```
-- .....(p.....L-],  
22
```

```
1 dropout_range = [0.1,0.2,0.3,0.4,0.5]  
2 kernels = [2,3]  
3 featureK = []  
4 featureD = []  
5 # Flatten the images into vectors (1D) for feed forward network  
6 train_images_3d = train_image_array.reshape((len(train_image_array),48,48,1))  
7 test_images_3d = test_image_array.reshape((len(test_image_array), 48,48,1))  
8 valid_images_3d = valid_image_array.reshape((len(valid_image_array), 48,48,1))  
9 results = []  
10 for i in dropout_range:  
11     for j in kernels:  
12         run_cnn_model(j,i)  
13         featureK.append(j)  
14         featureD.append(i)  
15 print(results)
```



Kernel Size: 2  
Dropout: 0.1  
Model: "sequential\_29"

Layer (type)	Output Shape	Param #
conv2d_113 (Conv2D)	(None, 16, 16, 64)	320
conv2d_114 (Conv2D)	(None, 15, 15, 64)	16448
max_pooling2d_57 (MaxPooling)	(None, 7, 7, 64)	0
dropout_57 (Dropout)	(None, 7, 7, 64)	0
conv2d_115 (Conv2D)	(None, 6, 6, 128)	32896
conv2d_116 (Conv2D)	(None, 5, 5, 128)	65664
max_pooling2d_58 (MaxPooling)	(None, 2, 2, 128)	0
dropout_58 (Dropout)	(None, 2, 2, 128)	0
flatten_29 (Flatten)	(None, 512)	0
dense_57 (Dense)	(None, 512)	262656
dense_58 (Dense)	(None, 7)	3591

Total params: 381,575  
Trainable params: 381,575  
Non-trainable params: 0

None

Epoch 1/10

28709/28709 [=====] - 4s 130us/step - loss: 3.7096 - acc: 0.

Epoch 2/10

28709/28709 [=====] - 1s 44us/step - loss: 1.9558 - acc: 0.2

Epoch 3/10

28709/28709 [=====] - 1s 44us/step - loss: 1.9367 - acc: 0.2

Epoch 4/10

28709/28709 [=====] - 1s 44us/step - loss: 1.9275 - acc: 0.2

Epoch 5/10

28709/28709 [=====] - 1s 44us/step - loss: 1.9317 - acc: 0.2

Epoch 6/10

28709/28709 [=====] - 1s 43us/step - loss: 1.9115 - acc: 0.2

Epoch 7/10

28709/28709 [=====] - 1s 43us/step - loss: 1.8791 - acc: 0.3

Epoch 8/10

28709/28709 [=====] - 1s 44us/step - loss: 1.8719 - acc: 0.3

Epoch 9/10

28709/28709 [=====] - 1s 44us/step - loss: 1.8369 - acc: 0.3

Epoch 10/10

28709/28709 [=====] - 1s 43us/step - loss: 1.8920 - acc: 0.3

3589/3589 [=====] - 1s 334us/step

Accuracy on Test samples: 0.24965171357339125

Kernel Size: 3

Dropout: 0.1

Model: "sequential\_30"

Layer (type)	Output Shape	Param #
=====	=====	=====

conv2d_117 (Conv2D)	(None, 16, 16, 64)	640
conv2d_118 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_59 (MaxPooling)	(None, 7, 7, 64)	0
dropout_59 (Dropout)	(None, 7, 7, 64)	0
conv2d_119 (Conv2D)	(None, 5, 5, 128)	73856
conv2d_120 (Conv2D)	(None, 3, 3, 128)	147584
max_pooling2d_60 (MaxPooling)	(None, 1, 1, 128)	0
dropout_60 (Dropout)	(None, 1, 1, 128)	0
flatten_30 (Flatten)	(None, 128)	0
dense_59 (Dense)	(None, 512)	66048
dense_60 (Dense)	(None, 7)	3591

=====  
Total params: 328,647  
Trainable params: 328,647  
Non-trainable params: 0

None

Epoch 1/10

28709/28709 [=====] - 4s 135us/step - loss: 4.3348 - acc: 0.

Epoch 2/10

28709/28709 [=====] - 1s 46us/step - loss: 4.2666 - acc: 0.2

Epoch 3/10

28709/28709 [=====] - 1s 46us/step - loss: 4.2487 - acc: 0.2

Epoch 4/10

28709/28709 [=====] - 1s 47us/step - loss: 4.2666 - acc: 0.2

Epoch 5/10

28709/28709 [=====] - 1s 46us/step - loss: 4.2508 - acc: 0.2

Epoch 6/10

28709/28709 [=====] - 1s 46us/step - loss: 4.2085 - acc: 0.3

Epoch 7/10

28709/28709 [=====] - 1s 46us/step - loss: 4.1103 - acc: 0.3

Epoch 8/10

28709/28709 [=====] - 1s 47us/step - loss: 4.0372 - acc: 0.3

Epoch 9/10

28709/28709 [=====] - 1s 47us/step - loss: 4.0335 - acc: 0.3

Epoch 10/10

28709/28709 [=====] - 1s 46us/step - loss: 4.0242 - acc: 0.3

3589/3589 [=====] - 1s 363us/step

Accuracy on Test samples: 0.3700195040567302

Kernel Size: 2

Dropout: 0.2

Model: "sequential\_31"

Layer (type)	Output Shape	Param #
conv2d_121 (Conv2D)	(None, 16, 16, 64)	320
conv2d_122 (Conv2D)	(None, 15, 15, 64)	16448
max_pooling2d_61 (MaxPooling)	(None, 7, 7, 64)	0
dropout_61 (Dropout)	(None, 7, 7, 64)	0

dropout_61 (Dropout)	(None, 7, 7, 32)	0
conv2d_123 (Conv2D)	(None, 6, 6, 128)	32896
conv2d_124 (Conv2D)	(None, 5, 5, 128)	65664
max_pooling2d_62 (MaxPooling)	(None, 2, 2, 128)	0
dropout_62 (Dropout)	(None, 2, 2, 128)	0
flatten_31 (Flatten)	(None, 512)	0
dense_61 (Dense)	(None, 512)	262656
dense_62 (Dense)	(None, 7)	3591
=====		
Total params: 381,575		
Trainable params: 381,575		
Non-trainable params: 0		

None

Epoch 1/10

28709/28709 [=====] - 4s 141us/step - loss: 5.6505 - acc: 0.

Epoch 2/10

28709/28709 [=====] - 1s 45us/step - loss: 5.5264 - acc: 0.2

Epoch 3/10

28709/28709 [=====] - 1s 45us/step - loss: 5.5252 - acc: 0.2

Epoch 4/10

28709/28709 [=====] - 1s 45us/step - loss: 5.5148 - acc: 0.2

Epoch 5/10

28709/28709 [=====] - 1s 45us/step - loss: 5.5090 - acc: 0.2

Epoch 6/10

28709/28709 [=====] - 1s 44us/step - loss: 5.5137 - acc: 0.2

Epoch 7/10

28709/28709 [=====] - 1s 44us/step - loss: 5.5278 - acc: 0.2

Epoch 8/10

28709/28709 [=====] - 1s 45us/step - loss: 5.5024 - acc: 0.2

Epoch 9/10

28709/28709 [=====] - 1s 45us/step - loss: 5.4945 - acc: 0.2

Epoch 10/10

28709/28709 [=====] - 1s 44us/step - loss: 5.5209 - acc: 0.2

3589/3589 [=====] - 1s 363us/step

Accuracy on Test samples: 0.2471440512719145

Kernel Size: 3

Dropout: 0.2

Model: "sequential\_32"

Layer (type)	Output Shape	Param #
=====		
conv2d_125 (Conv2D)	(None, 16, 16, 64)	640
conv2d_126 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_63 (MaxPooling)	(None, 7, 7, 64)	0
dropout_63 (Dropout)	(None, 7, 7, 64)	0
conv2d_127 (Conv2D)	(None, 5, 5, 128)	73856
conv2d_128 (Conv2D)	(None, 3, 3, 128)	147584
max_pooling2d_64 (MaxPooling)	(None, 1, 1, 128)	0



dropout_64 (Dropout)	(None, 1, 1, 128)	0
flatten_32 (Flatten)	(None, 128)	0
dense_63 (Dense)	(None, 512)	66048
dense_64 (Dense)	(None, 7)	3591
=====		
Total params: 328,647		
Trainable params: 328,647		
Non-trainable params: 0		

None

Epoch 1/10

28709/28709 [=====] - 4s 145us/step - loss: 2.1478 - acc: 0.

Epoch 2/10

28709/28709 [=====] - 1s 46us/step - loss: 1.9520 - acc: 0.2

Epoch 3/10

28709/28709 [=====] - 1s 46us/step - loss: 1.9378 - acc: 0.2

Epoch 4/10

28709/28709 [=====] - 1s 47us/step - loss: 1.9059 - acc: 0.2

Epoch 5/10

28709/28709 [=====] - 1s 47us/step - loss: 1.8848 - acc: 0.3

Epoch 6/10

28709/28709 [=====] - 1s 47us/step - loss: 1.8792 - acc: 0.3

Epoch 7/10

28709/28709 [=====] - 1s 46us/step - loss: 1.9025 - acc: 0.2

Epoch 8/10

28709/28709 [=====] - 1s 46us/step - loss: 1.9804 - acc: 0.2

Epoch 9/10

28709/28709 [=====] - 1s 46us/step - loss: 1.9621 - acc: 0.2

Epoch 10/10

28709/28709 [=====] - 1s 47us/step - loss: 1.9537 - acc: 0.2

3589/3589 [=====] - 1s 399us/step

Accuracy on Test samples: 0.24965171357339125

Kernel Size: 2

Dropout: 0.3

Model: "sequential\_33"

Layer (type)	Output Shape	Param #
=====		
conv2d_129 (Conv2D)	(None, 16, 16, 64)	320
conv2d_130 (Conv2D)	(None, 15, 15, 64)	16448
max_pooling2d_65 (MaxPooling)	(None, 7, 7, 64)	0
dropout_65 (Dropout)	(None, 7, 7, 64)	0
conv2d_131 (Conv2D)	(None, 6, 6, 128)	32896
conv2d_132 (Conv2D)	(None, 5, 5, 128)	65664
max_pooling2d_66 (MaxPooling)	(None, 2, 2, 128)	0
dropout_66 (Dropout)	(None, 2, 2, 128)	0
flatten_33 (Flatten)	(None, 512)	0
dense_65 (Dense)	(None, 512)	262656

dense_66 (Dense)	(None, 7)	3591
------------------	-----------	------

=====

Total params: 381,575  
Trainable params: 381,575  
Non-trainable params: 0

---

None

Epoch 1/10  
28709/28709 [=====] - 4s 146us/step - loss: 2.0020 - acc: 0.  
Epoch 2/10  
28709/28709 [=====] - 1s 45us/step - loss: 1.9467 - acc: 0.2  
Epoch 3/10  
28709/28709 [=====] - 1s 45us/step - loss: 1.9280 - acc: 0.2  
Epoch 4/10  
28709/28709 [=====] - 1s 45us/step - loss: 1.9140 - acc: 0.2  
Epoch 5/10  
28709/28709 [=====] - 1s 45us/step - loss: 1.8755 - acc: 0.3  
Epoch 6/10  
28709/28709 [=====] - 1s 45us/step - loss: 1.8854 - acc: 0.3  
Epoch 7/10  
28709/28709 [=====] - 1s 45us/step - loss: 1.8927 - acc: 0.2  
Epoch 8/10  
28709/28709 [=====] - 1s 45us/step - loss: 1.8536 - acc: 0.3  
Epoch 9/10  
28709/28709 [=====] - 1s 45us/step - loss: 1.8521 - acc: 0.3  
Epoch 10/10  
28709/28709 [=====] - 1s 45us/step - loss: 1.8289 - acc: 0.3  
3589/3589 [=====] - 1s 407us/step  
Accuracy on Test samples: 0.3382557815630544  
Kernel Size: 3  
Dropout: 0.3  
Model: "sequential\_34"

Layer (type)	Output Shape	Param #
conv2d_133 (Conv2D)	(None, 16, 16, 64)	640
conv2d_134 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_67 (MaxPooling)	(None, 7, 7, 64)	0
dropout_67 (Dropout)	(None, 7, 7, 64)	0
conv2d_135 (Conv2D)	(None, 5, 5, 128)	73856
conv2d_136 (Conv2D)	(None, 3, 3, 128)	147584
max_pooling2d_68 (MaxPooling)	(None, 1, 1, 128)	0
dropout_68 (Dropout)	(None, 1, 1, 128)	0
flatten_34 (Flatten)	(None, 128)	0
dense_67 (Dense)	(None, 512)	66048
dense_68 (Dense)	(None, 7)	3591

=====

Total params: 328,647  
Trainable params: 328,647  
Non-trainable params: 0

---

None  
Epoch 1/10  
28709/28709 [=====] - 4s 156us/step - loss: 1.9843 - acc: 0.  
Epoch 2/10  
28709/28709 [=====] - 1s 47us/step - loss: 1.9349 - acc: 0.2  
Epoch 3/10  
28709/28709 [=====] - 1s 48us/step - loss: 1.8743 - acc: 0.3  
Epoch 4/10  
28709/28709 [=====] - 1s 47us/step - loss: 1.8782 - acc: 0.3  
Epoch 5/10  
28709/28709 [=====] - 1s 47us/step - loss: 1.8071 - acc: 0.3  
Epoch 6/10  
28709/28709 [=====] - 1s 46us/step - loss: 1.7780 - acc: 0.3  
Epoch 7/10  
28709/28709 [=====] - 1s 47us/step - loss: 1.7639 - acc: 0.3  
Epoch 8/10  
28709/28709 [=====] - 1s 46us/step - loss: 1.7149 - acc: 0.4  
Epoch 9/10  
28709/28709 [=====] - 1s 46us/step - loss: 1.6695 - acc: 0.4  
Epoch 10/10  
28709/28709 [=====] - 1s 46us/step - loss: 1.6500 - acc: 0.4  
3589/3589 [=====] - 2s 420us/step  
Accuracy on Test samples: 0.45082195598489894  
Kernel Size: 2  
Dropout: 0.4  
Model: "sequential\_35"

Layer (type)	Output Shape	Param #
conv2d_137 (Conv2D)	(None, 16, 16, 64)	320
conv2d_138 (Conv2D)	(None, 15, 15, 64)	16448
max_pooling2d_69 (MaxPooling)	(None, 7, 7, 64)	0
dropout_69 (Dropout)	(None, 7, 7, 64)	0
conv2d_139 (Conv2D)	(None, 6, 6, 128)	32896
conv2d_140 (Conv2D)	(None, 5, 5, 128)	65664
max_pooling2d_70 (MaxPooling)	(None, 2, 2, 128)	0
dropout_70 (Dropout)	(None, 2, 2, 128)	0
flatten_35 (Flatten)	(None, 512)	0
dense_69 (Dense)	(None, 512)	262656
dense_70 (Dense)	(None, 7)	3591

=====  
Total params: 381,575  
Trainable params: 381,575  
Non-trainable params: 0

None  
Epoch 1/10  
28709/28709 [=====] - 4s 157us/step - loss: 3.9552 - acc: 0.  
Epoch 2/10  
28709/28709 [=====] - 1s 46us/step - loss: 3.8373 - acc: 0.2  
Epoch 3/10  
28709/28709 [=====] - 1s 45us/step - loss: 3.8247 - acc: 0.2

Epoch 4/10  
 28709/28709 [=====] - 1s 45us/step - loss: 3.8246 - acc: 0.2  
 Epoch 5/10  
 28709/28709 [=====] - 1s 45us/step - loss: 3.8046 - acc: 0.2  
 Epoch 6/10  
 28709/28709 [=====] - 1s 44us/step - loss: 3.7788 - acc: 0.3  
 Epoch 7/10  
 28709/28709 [=====] - 1s 45us/step - loss: 3.7734 - acc: 0.3  
 Epoch 8/10  
 28709/28709 [=====] - 1s 45us/step - loss: 3.7678 - acc: 0.3  
 Epoch 9/10  
 28709/28709 [=====] - 1s 45us/step - loss: 3.7833 - acc: 0.3  
 Epoch 10/10  
 28709/28709 [=====] - 1s 46us/step - loss: 3.7560 - acc: 0.3  
 3589/3589 [=====] - 2s 445us/step  
 Accuracy on Test samples: 0.34717191419052723  
 Kernel Size: 3  
 Dropout: 0.4  
 Model: "sequential\_36"

Layer (type)	Output Shape	Param #
conv2d_141 (Conv2D)	(None, 16, 16, 64)	640
conv2d_142 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_71 (MaxPooling)	(None, 7, 7, 64)	0
dropout_71 (Dropout)	(None, 7, 7, 64)	0
conv2d_143 (Conv2D)	(None, 5, 5, 128)	73856
conv2d_144 (Conv2D)	(None, 3, 3, 128)	147584
max_pooling2d_72 (MaxPooling)	(None, 1, 1, 128)	0
dropout_72 (Dropout)	(None, 1, 1, 128)	0
flatten_36 (Flatten)	(None, 128)	0
dense_71 (Dense)	(None, 512)	66048
dense_72 (Dense)	(None, 7)	3591

=====

Total params: 328,647  
 Trainable params: 328,647  
 Non-trainable params: 0

None

Epoch 1/10  
 28709/28709 [=====] - 5s 162us/step - loss: 2.0823 - acc: 0.  
 Epoch 2/10  
 28709/28709 [=====] - 1s 46us/step - loss: 1.9556 - acc: 0.2  
 Epoch 3/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9332 - acc: 0.2  
 Epoch 4/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.8788 - acc: 0.3  
 Epoch 5/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.8616 - acc: 0.3  
 Epoch 6/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9176 - acc: 0.2

Epoch 7/10  
 28709/28709 [=====] - 1s 46us/step - loss: 1.8957 - acc: 0.3  
 Epoch 8/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.8530 - acc: 0.3  
 Epoch 9/10  
 28709/28709 [=====] - 1s 46us/step - loss: 1.8071 - acc: 0.3  
 Epoch 10/10  
 28709/28709 [=====] - 1s 46us/step - loss: 1.7736 - acc: 0.3  
 3589/3589 [=====] - 2s 444us/step  
 Accuracy on Test samples: 0.41432153804118205  
 Kernel Size: 2  
 Dropout: 0.5  
 Model: "sequential\_37"

Layer (type)	Output Shape	Param #
=====		
conv2d_145 (Conv2D)	(None, 16, 16, 64)	320
conv2d_146 (Conv2D)	(None, 15, 15, 64)	16448
max_pooling2d_73 (MaxPooling)	(None, 7, 7, 64)	0
dropout_73 (Dropout)	(None, 7, 7, 64)	0
conv2d_147 (Conv2D)	(None, 6, 6, 128)	32896
conv2d_148 (Conv2D)	(None, 5, 5, 128)	65664
max_pooling2d_74 (MaxPooling)	(None, 2, 2, 128)	0
dropout_74 (Dropout)	(None, 2, 2, 128)	0
flatten_37 (Flatten)	(None, 512)	0
dense_73 (Dense)	(None, 512)	262656
dense_74 (Dense)	(None, 7)	3591
=====		

Total params: 381,575  
 Trainable params: 381,575  
 Non-trainable params: 0

None  
 Epoch 1/10  
 28709/28709 [=====] - 5s 169us/step - loss: 2.0850 - acc: 0.  
 Epoch 2/10  
 28709/28709 [=====] - 1s 48us/step - loss: 1.9683 - acc: 0.2  
 Epoch 3/10  
 28709/28709 [=====] - 1s 48us/step - loss: 1.9480 - acc: 0.2  
 Epoch 4/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9409 - acc: 0.2  
 Epoch 5/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9328 - acc: 0.2  
 Epoch 6/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9206 - acc: 0.2  
 Epoch 7/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.8985 - acc: 0.2  
 Epoch 8/10  
 28709/28709 [=====] - 1s 46us/step - loss: 1.8919 - acc: 0.3  
 Epoch 9/10  
 28709/28709 [=====] - 1s 46us/step - loss: 1.8900 - acc: 0.3  
 Epoch 10/10

28709/28709 [=====] - 1s 46us/step - loss: 1.8610 - acc: 0.3  
 3589/3589 [=====] - 2s 465us/step  
 Accuracy on Test samples: 0.35775982171056264  
 Kernel Size: 3  
 Dropout: 0.5  
 Model: "sequential\_38"

Layer (type)	Output Shape	Param #
conv2d_149 (Conv2D)	(None, 16, 16, 64)	640
conv2d_150 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_75 (MaxPooling)	(None, 7, 7, 64)	0
dropout_75 (Dropout)	(None, 7, 7, 64)	0
conv2d_151 (Conv2D)	(None, 5, 5, 128)	73856
conv2d_152 (Conv2D)	(None, 3, 3, 128)	147584
max_pooling2d_76 (MaxPooling)	(None, 1, 1, 128)	0
dropout_76 (Dropout)	(None, 1, 1, 128)	0
flatten_38 (Flatten)	(None, 128)	0
dense_75 (Dense)	(None, 512)	66048
dense_76 (Dense)	(None, 7)	3591

=====

Total params: 328,647  
 Trainable params: 328,647  
 Non-trainable params: 0

None

Epoch 1/10  
 28709/28709 [=====] - 5s 174us/step - loss: 2.1674 - acc: 0.  
 Epoch 2/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9582 - acc: 0.2  
 Epoch 3/10  
 28709/28709 [=====] - 1s 48us/step - loss: 1.9515 - acc: 0.2  
 Epoch 4/10  
 28709/28709 [=====] - 1s 48us/step - loss: 1.9384 - acc: 0.2  
 Epoch 5/10  
 28709/28709 [=====] - 1s 46us/step - loss: 1.9252 - acc: 0.2  
 Epoch 6/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9082 - acc: 0.2  
 Epoch 7/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.8795 - acc: 0.3  
 Epoch 8/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.8527 - acc: 0.3  
 Epoch 9/10  
 28709/28709 [=====] - 1s 48us/step - loss: 1.9342 - acc: 0.2  
 Epoch 10/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9389 - acc: 0.2  
 3589/3589 [=====] - 2s 493us/step  
 Accuracy on Test samples: 0.289774310397019  
 [0.24965171357339125, 0.3700195040567302, 0.2471440512719145, 0.24965171357339125, 0.

```

1 best=(np.argmax(results))
2 print('Best Feature Pair is: {0} and {1}'.format(featureK[best], featureD[best]))
3 print("Best performance from training:",results[best])
4 kSize = featureK[best]
5 dropout_rate = featureD[best]
6 common_features = [Conv2D(64, kernel_size=(kSize), strides=3, activation='relu', input_
7                      Conv2D(64, kernel_size=(kSize), activation='relu'),
8                      MaxPooling2D(pool_size=(2,2)), Dropout(dropout_rate),
9                      Conv2D(128, kernel_size=(kSize), activation='relu'),
10                     Conv2D(128, kernel_size=(kSize), activation='relu'),
11                     MaxPooling2D(pool_size=(2,2)), Dropout(dropout_rate), Flatten(),]
12
13 classifier = [Dense(512, activation='relu'), Dense(7, activation='relu'),]
14 cnn_model = Sequential(common_features+classifier)
15 print("Kernel Size: {0} \nDropout: {1}".format(kSize,dropout_rate))
16 print(cnn_model.summary())
17 cnn_model.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['accuracy'
18 # Train model
19 cnn_model.fit(train_images_3d, to_categorical(train_labels), epochs=10, batch_size=256,
20
21 performance = cnn_model.evaluate(test_images_3d, to_categorical(test_labels))
22 print("Accuracy on Test samples: {0}".format(performance[1]))

```



Best Feature Pair is: 3 and 0.3  
 Best performance from training: 0.45082195598489894  
 Kernel Size: 3  
 Dropout: 0.3  
 Model: "sequential\_40"

Layer (type)	Output Shape	Param #
conv2d_157 (Conv2D)	(None, 16, 16, 64)	640
conv2d_158 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_79 (MaxPooling)	(None, 7, 7, 64)	0
dropout_79 (Dropout)	(None, 7, 7, 64)	0
conv2d_159 (Conv2D)	(None, 5, 5, 128)	73856
conv2d_160 (Conv2D)	(None, 3, 3, 128)	147584
max_pooling2d_80 (MaxPooling)	(None, 1, 1, 128)	0
dropout_80 (Dropout)	(None, 1, 1, 128)	0
flatten_40 (Flatten)	(None, 128)	0
dense_79 (Dense)	(None, 512)	66048
dense_80 (Dense)	(None, 7)	3591

=====  
 Total params: 328,647  
 Trainable params: 328,647  
 Non-trainable params: 0

None  
 Epoch 1/10  
 28709/28709 [=====] - 5s 182us/step - loss: 1.9924 - acc: 0.  
 Epoch 2/10  
 28709/28709 [=====] - 1s 48us/step - loss: 1.9631 - acc: 0.2  
 Epoch 3/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9470 - acc: 0.2  
 Epoch 4/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9074 - acc: 0.2  
 Epoch 5/10  
 28709/28709 [=====] - 1s 48us/step - loss: 1.8598 - acc: 0.3  
 Epoch 6/10  
 28709/28709 [=====] - 1s 48us/step - loss: 1.8313 - acc: 0.3  
 Epoch 7/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9253 - acc: 0.2  
 Epoch 8/10  
 28709/28709 [=====] - 1s 47us/step - loss: 1.9469 - acc: 0.2  
 Epoch 9/10  
 28709/28709 [=====] - 1s 49us/step - loss: 1.8832 - acc: 0.3  
 Epoch 10/10  
 28709/28709 [=====] - 1s 48us/step - loss: 1.8336 - acc: 0.3  
 3589/3589 [=====] - 2s 522us/step  
 Accuracy on Test samples: 0.2449150181191982



```

1 def optimize_cnn(hyperparameter):
2
3     # Define model using hyperparameters
4     cnn_model = Sequential([Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], ac
5                               Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], activation='relu'
6                               MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
7                               Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu'
8                               Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu'
9                               MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
10                               Flatten(),
11                               Dense(512, activation='relu'),
12                               Dense(7, activation='softmax'),])
13
14     cnn_model.compile(optimizer=hyperparameter['optimizer'], loss='categorical_crossentropy',
15                       time_callback = TimeHistory())
16     hist = cnn_model.fit(train_images_3d, to_categorical(train_labels), epochs=10, batch_size=10,
17                         # print(hist.history.accuracy)
18                         print("Total Training Time is (s): ", sum(time_callback.times))
19                         # Evaluate accuracy on validation data
20                         performance = cnn_model.evaluate(valid_images_3d, to_categorical(valid_labels), verbose=0))
21
22     print("Hyperparameters: ", hyperparameter, "Accuracy: ", performance[1])
23     print("-----")
24     # We want to minimize loss i.e. negative of accuracy
25     return({"status": STATUS_OK, "loss": -1*performance[1], "model":cnn_model})
26
27
28 # Define search space for hyper-parameters
29 space = {
30     # The kernel_size for convolutions:
31     'conv_kernel_size': hp.choice('conv_kernel_size', [1, 3, 5]),
32     # Uniform distribution in finding appropriate dropout values
33     'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.35),
34     # Choice of optimizer
35     'optimizer': hp.choice('optimizer', ['Adam', 'sgd']),
36 }
37
38 trials = Trials()
39
40 # Find the best hyperparameters
41 best = fmin(
42     optimize_cnn,
43     space,
44     algo=tpe.suggest,
45     trials=trials,
46     max_evals=25,
47 )
48
49 print("=====")
50 print("Best Hyperparameters", best)
51
52 # You can retrain the final model with optimal hyperparameters on train+validation data
53
54 # Or you can use the model returned directly
55 # Find trial which has minimum loss value and use that model to perform evaluation on test data

```

```
55 # Find trial which has minimum loss value and use that model to perform evaluation on test set
56 test_model = trials.results[np.argmin([r['loss'] for r in trials.results])]['model']
57
58 performance = test_model.evaluate(test_images_3d, to_categorical(test_labels))
59
60 print("=====")
61 print("Test Accuracy: ", performance[1])
```



```
Total Training Time is (s):
28.2460458278656
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.14668735307891242, 'optimizer': 'sgd'}
Accuracy:
0.3399275564265537
-----
Total Training Time is (s):
23.3261399269104
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.17386172811548062, 'optimizer': 'sgd'}
Accuracy:
0.3265533575144077
-----
Total Training Time is (s):
29.898162126541138
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.16029133971157797, 'optimizer': 'Adam'}
Accuracy:
0.5770409585174726
-----
Total Training Time is (s):
28.340755224227905
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.19075674893547134, 'optimizer': 'sgd'}
Accuracy:
0.3558094176692397
-----
Total Training Time is (s):
28.23865532875061
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.2915313031622959, 'optimizer': 'sgd'}
Accuracy:
0.3349122318236002
-----
Total Training Time is (s):
29.37062120437622
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.34814906012233493, 'optimizer': 'Adam'}
Accuracy:
0.5884647534464222
-----
Total Training Time is (s):
29.427371740341187
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.12730390537713165, 'optimizer': 'Adam'}
Accuracy:
0.5795486208189494
-----
Total Training Time is (s):
23.43778681755066
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.2914162541212033, 'optimizer': 'sgd'}
Accuracy:
0.32683198663410484
-----
Total Training Time is (s):
23.60374140739441
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.23456208751505514, 'optimizer': 'sgd'}
Accuracy:
```

0.3413207021786596

-----  
Total Training Time is (s):

25.58343744277954

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.21816456764498765, 'optimizer': 'Adam'}

Accuracy:

0.4611312343312625

-----  
40%|██████ | 10/25 [04:58<07:27, 29.83s/it, best loss: -0.5884647534464222]

-----  
KeyboardInterrupt Traceback (most recent call last)

<ipython-input-18-dd4c7d03b67d> in <module>()

```
    44         algo=tpe.suggest,
    45         trials=trials,
--> 46         max_evals=25,
    47     )
    48
```

-----  
⌵ 12 frames -----

/tensorflow-1.15.0/python3.6/tensorflow\_core/python/client/session.py in \_\_call\_\_(self,

```
    1470         ret = tf_session.TF_SessionRunCallable(self._session._session,
    1471                                                  self._handle, args,
-> 1472                                                  run_metadata_ptr)
    1473         if run_metadata:
    1474             proto_data = tf_session.TF_GetBuffer(run_metadata_ptr)
```

KeyboardInterrupt:

SEARCH STACK OVERFLOW

## ▼ FINE TUNING

```
1 from keras.datasets import mnist
2 from keras.models import Model
3 from keras.layers import Input
4 #Load MNIST dataset
5 (train_imagesf, train_labelsf), (test_imagesf, test_labelsf) = mnist.load_data()
6 train_imagesf = (train_imagesf / 255) - 0.5
7 test_imagesf = (test_imagesf / 255) - 0.5
```

☞ Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 0s 0us/step

```
1 train_images_3df = train_imagesf.reshape(60000,28,28,1)
2 test_images_3df = test_imagesf.reshape(10000,28,28,1)
3 train_images_3df = train_images_3df[:28709]
4 train_labelsf = train_labelsf[:28709]
5 test_images_3df = test_images_3df[:3589]
6 test_labelsf = test_labelsf[:3589]
7
```

```

8 Base_feature_model = Sequential([Conv2D(32, kernel_size=3, activation='relu'),
9     Conv2D(32, kernel_size=3, activation='relu'),
10    MaxPooling2D(pool_size=(2,2)), Dropout(0.25),
11    Conv2D(64, kernel_size=3, activation='relu'),
12    Conv2D(64, kernel_size=3, activation='relu'),
13    MaxPooling2D(pool_size=(2,2)), Dropout(0.25),
14    Dense(512, activation='relu'), Flatten(),])
15
16 # Define task-specific classification layers
17 Classifier_mnist = Dense(10, activation='softmax', name='MNIST')
18 Classifier_fer = Dense(7, activation='softmax', name='FER')
19
20 # Instantiate a Tensor to feed Input (Input Layer)
21 mnist_input = Input(shape=(28,28,1))
22 fer_input = Input(shape=(48,48,1))
23
24 # Call Base_feature_model over the mnist images
25 mnist_features = Base_feature_model(mnist_input)
26
27 # Call Base_feature_model over the fashion-mnist images
28 fer_features = Base_feature_model(fer_input)
29
30 # Call mnist_prediction layer over the mnist images
31 # mnist_prediction represents the predicted output for mnist dataset
32 mnist_prediction = Classifier_mnist(mnist_features)
33
34 # Call fashion_mnist_prediction layer over the mnist images
35 # fashion_mnist_prediction represents the predicted output for fashion-mnist dataset
36 fer_prediction = Classifier_fer(fer_features)
37
38
39 joint_model = Model(inputs=[mnist_input, fer_input],
40     outputs=[mnist_prediction, fer_prediction])
41
42 print(joint_model.summary())
43
44 joint_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'],)
45
46 joint_model.fit([train_images_3df, train_images_3d],
47     [to_categorical(train_labelsf), to_categorical(train_labels)],
48     epochs=2, batch_size=1024,)
49 performance = joint_model.evaluate([test_images_3df, valid_images_3d],
50     [to_categorical(test_labelsf),
51     to_categorical(valid_labels)], verbose=1)
52
53 print("===\nMNIST Accuracy: {0}\nFER Accuracy: {1}".format(performance[3], performance[
54 performance = joint_model.evaluate([test_images_3df, test_images_3d],
55     [to_categorical(test_labelsf),
56     to_categorical(test_labels)], verbose=1)
57
58 print("===\nMNIST Accuracy: {0}\nFER Accuracy: {1}".format(performance[3], performance[

```



Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 28, 28, 1)	0	
input_2 (InputLayer)	(None, 48, 48, 1)	0	
sequential_2 (Sequential)	multiple	98272	input_1[0][0] input_2[0][0]
MNIST (Dense)	(None, 10)	81930	sequential_2[1][0]
FER (Dense)	(None, 7)	290311	sequential_2[2][0]
Total params: 470,513			
Trainable params: 470,513			
Non-trainable params: 0			

None

WARNING:tensorflow:From /tensorflow-1.15.0/python3.6/tensorflow\_core/python/ops/nn\_in  
Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

Epoch 1/2

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

28709/28709 [=====] - 18s 611us/step - loss: 0.5884 - MNIST\_  
Epoch 2/2

28709/28709 [=====] - 3s 116us/step - loss: 0.4026 - MNIST\_  
3589/3589 [=====] - 1s 209us/step

===

MNIST Accuracy: 0.9875731395127282

FER Accuracy: 0.8647852702936644

3589/3589 [=====] - 1s 145us/step

===

MNIST Accuracy: 0.9875731395127282

FER Accuracy: 0.8625562360739103

## ▼ Fine tuning using the MNIST Model.

```

1 import cv2
2 from skimage.transform import resize
3 # Resizing the images.
4 flatten_train_images = train_image_array.reshape((-1, 48,48))
5 flatten_test_images = test_image_array.reshape((-1, 48,48))
6 flatten_valid_images = valid_image_array.reshape((-1,48,48))
7 resized_images_train = []
8 resized_images_test = []
9 resized_images_valid = []
10 # for i in range(len(flatten_test_images)):
11 for i in range(len(train_pixels)):
12     img = flatten_train_images[i]
13     img=cv2.resize(img, dsize=(28,28),interpolation=cv2.INTER_CUBIC)
14     resized_images_train.append(img)
15 resized_images_train = np.array(resized_images_train)
16
17 for i in range(len(valid_pixels)):
18     img = flatten_valid_images[i]
19     img=cv2.resize(img, dsize=(28,28),interpolation=cv2.INTER_CUBIC)
20     resized_images_valid.append(img)
21 resized_images_valid = np.array(resized_images_valid)
22
23
24 for i in range(len(test_pixels)):
25     img = flatten_test_images[i]
26     img=cv2.resize(img, dsize=(28,28),interpolation=cv2.INTER_CUBIC)
27     resized_images_test.append(img)
28 resized_images_test = np.array(resized_images_test)
29

```

```

1 from keras.layers import Conv2D, Flatten, MaxPooling2D
2
3 # Define 2 groups of layers: features layer (convolutions) and classification layer
4 common_features = [Conv2D(32, kernel_size=3, activation='relu', input_shape=(28,28,1)),
5                     Conv2D(32, kernel_size=3, activation='relu'),
6                     MaxPooling2D(pool_size=(2,2)), Dropout(0.25),
7                     Conv2D(64, kernel_size=3, activation='relu'),
8                     Conv2D(64, kernel_size=3, activation='relu'),
9                     MaxPooling2D(pool_size=(2,2)), Dropout(0.25), Flatten(),]
10 classifier = [Dense(512, activation='relu'), Dense(10, activation='softmax'),]
11
12 cnn_model = Sequential(common_features+classifier)
13
14 print(cnn_model.summary()) # Compare number of parameteres against FFN
15 cnn_model.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['accuracy'])
16
17 train_images_3df = train_imagesf.reshape(60000,28,28,1)
18 test_images_3df = test_imagesf.reshape(10000,28,28,1)
19
20 cnn_model.fit(train_images_3df, to_categorical(train_labelsf), epochs=10, batch_size=256)
21 performance = cnn_model.evaluate(test_images_3df, to_categorical(test_labelsf))
22
23 print("Accuracy on Test samples: {}".format(performance[1]))

```

Model: "sequential\_29"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 26, 26, 32)	320
conv2d_10 (Conv2D)	(None, 24, 24, 32)	9248
max_pooling2d_5 (MaxPooling2D)	(None, 12, 12, 32)	0
dropout_31 (Dropout)	(None, 12, 12, 32)	0
conv2d_11 (Conv2D)	(None, 10, 10, 64)	18496
conv2d_12 (Conv2D)	(None, 8, 8, 64)	36928
max_pooling2d_6 (MaxPooling2D)	(None, 4, 4, 64)	0
dropout_32 (Dropout)	(None, 4, 4, 64)	0
flatten_3 (Flatten)	(None, 1024)	0
dense_31 (Dense)	(None, 512)	524800
dense_32 (Dense)	(None, 10)	5130
Total params: 594,922		
Trainable params: 594,922		
Non-trainable params: 0		

None

Epoch 1/10

60000/60000 [=====] - 9s 145us/step - loss: 0.2772 - acc: 0.9

Epoch 2/10

60000/60000 [=====] - 2s 39us/step - loss: 0.0598 - acc: 0.9

Epoch 3/10

60000/60000 [=====] - 2s 39us/step - loss: 0.0429 - acc: 0.9

Epoch 4/10

60000/60000 [=====] - 2s 39us/step - loss: 0.0336 - acc: 0.9

Epoch 5/10

60000/60000 [=====] - 2s 39us/step - loss: 0.0272 - acc: 0.9

Epoch 6/10

60000/60000 [=====] - 2s 39us/step - loss: 0.0242 - acc: 0.9

Epoch 7/10

60000/60000 [=====] - 2s 39us/step - loss: 0.0209 - acc: 0.9

Epoch 8/10

60000/60000 [=====] - 2s 38us/step - loss: 0.0188 - acc: 0.9

Epoch 9/10

60000/60000 [=====] - 2s 38us/step - loss: 0.0160 - acc: 0.9

Epoch 10/10

60000/60000 [=====] - 2s 39us/step - loss: 0.0148 - acc: 0.9

10000/10000 [=====] - 1s 124us/step

Accuracy on Test samples: 0.9944

1 for l in common\_features:

2 l.trainable = False

3



```
4 classifier2 = [Dense(512, activation='relu'), Dense(7, activation='softmax'),]
5
6 cnn_model.add(Dense(7,activation='softmax'))
7 print(cnn_model.summary())
8 resized_images_valid_3d = resized_images_valid.reshape(3589,28,28,1)
9 resized_images_train_3d = resized_images_train.reshape(28709,28,28,1)
10 cnn_model.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['accuracy']
11
12 cnn_model.fit(resized_images_train_3d, to_categorical(train_labels), epochs=8, batch_si
13 performance = cnn_model.evaluate(resized_images_valid_3d, to_categorical(valid_labels))
14
15 print("Accuracy on Test samples: {0}".format(performance[1]))
```



Model: "sequential\_29"

Layer (type)	Output Shape	Param #
=====		
conv2d_9 (Conv2D)	(None, 26, 26, 32)	320
conv2d_10 (Conv2D)	(None, 24, 24, 32)	9248
max_pooling2d_5 (MaxPooling2D)	(None, 12, 12, 32)	0
dropout_31 (Dropout)	(None, 12, 12, 32)	0
conv2d_11 (Conv2D)	(None, 10, 10, 64)	18496
conv2d_12 (Conv2D)	(None, 8, 8, 64)	36928
max_pooling2d_6 (MaxPooling2D)	(None, 4, 4, 64)	0
dropout_32 (Dropout)	(None, 4, 4, 64)	0
flatten_3 (Flatten)	(None, 1024)	0
dense_31 (Dense)	(None, 512)	524800
dense_32 (Dense)	(None, 10)	5130
dense_35 (Dense)	(None, 7)	77
=====		

Total params: 659,914

Trainable params: 594,922

Non-trainable params: 64,992

None

Epoch 1/8

28709/28709 [=====] - 2s 76us/step - loss: 1.8593 - acc: 0.2

Epoch 2/8

28709/28709 [=====] - 1s 23us/step - loss: 1.7825 - acc: 0.2

Epoch 3/8

28709/28709 [=====] - 1s 25us/step - loss: 1.7408 - acc: 0.3

Epoch 4/8

28709/28709 [=====] - 1s 25us/step - loss: 1.7033 - acc: 0.3

Epoch 5/8

28709/28709 [=====] - 1s 23us/step - loss: 1.6698 - acc: 0.3

Epoch 6/8

28709/28709 [=====] - 1s 24us/step - loss: 1.6397 - acc: 0.3

Epoch 7/8

28709/28709 [=====] - 1s 24us/step - loss: 1.6091 - acc: 0.3

Epoch 8/8

28709/28709 [=====] - 1s 24us/step - loss: 1.5825 - acc: 0.4

3589/3589 [=====] - 1s 232us/step

Accuracy on Test samples: 0.4093062134299248

## ▼ Data Augmentation

```
1 from keras.preprocessing.image import ImageDataGenerator
```

```

2 from matplotlib import pyplot
3 datagen = ImageDataGenerator()
4
5 #feature standardization
6 x_train = (train_image_array)
7 x_valid = (valid_image_array)
8 x_test = (test_image_array)
9
10 x_train = x_train.reshape((x_train.shape[0], 48, 48, 1))
11 x_valid = x_valid.reshape((x_valid.shape[0], 48, 48, 1))
12 x_test = x_test.reshape((x_test.shape[0], 48, 48, 1))
13
14 x_train = x_train.astype('float32')
15 x_valid = x_valid.astype('float32')
16 x_test = x_test.astype('float32')
17
18 datagen = ImageDataGenerator(featurewise_center=True, featurewise_std_normalization=True)
19 datagen.fit(x_train)
20 datagen.fit(x_valid)
21 datagen.fit(x_test)


1 def optimize_cnn(hyperparameter):
2
3     # Define model using hyperparameters
4     cnn_model = Sequential([Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], activation='relu'),
5                             Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], activation='relu'),
6                             MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
7                             Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu'),
8                             Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu'),
9                             MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
10                             Flatten(),
11                             Dense(512, activation='relu'),
12                             Dense(7, activation='softmax'),])
13
14     cnn_model.compile(optimizer=hyperparameter['optimizer'], loss='categorical_crossentropy')
15     time_callback = TimeHistory()
16     hist = cnn_model.fit(x_train, to_categorical(train_labels), epochs=10, batch_size=256)
17     # print(hist.history.accuracy)
18     print("Total Training Time is (s): ", sum(time_callback.times))
19     # Evaluate accuracy on validation data
20     performance = cnn_model.evaluate(x_valid, to_categorical(valid_labels), verbose=0)
21
22     print("Hyperparameters: ", hyperparameter, "Accuracy: ", performance[1])
23     print("-----")
24     # We want to minimize loss i.e. negative of accuracy
25     return({"status": STATUS_OK, "loss": -1*performance[1], "model":cnn_model})
26
27
28 # Define search space for hyper-parameters
29 space = {
30     # The kernel_size for convolutions:
31     'conv_kernel_size': hp.choice('conv_kernel_size', [1, 3, 5]),
32     # Uniform distribution in finding appropriate dropout values
33     'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.35),
34     # Choice of optimizer

```

```

34     # CHOICE OF OPTIMIZER
35     'optimizer': hp.choice('optimizer', ['Adam', 'sgd']),
36 }
37
38 trials = Trials()
39
40 # Find the best hyperparameters
41 best = fmin(
42     optimize_cnn,
43     space,
44     algo=tpe.suggest,
45     trials=trials,
46     max_evals=25,
47 )
48
49 print("=====")
50 print("Best Hyperparameters", best)
51
52 # You can retrain the final model with optimal hyperparameters on train+validation data
53
54 # Or you can use the model returned directly
55 # Find trial which has minimum loss value and use that model to perform evaluation on test data
56 test_model = trials.results[np.argmin([r['loss'] for r in trials.results])]['model']
57
58 performance = test_model.evaluate(x_test, to_categorical(test_labels))
59
60 print("=====")
61 print("Test Accuracy: ", performance[1])

```



```
Total Training Time is (s):
26.69492268562317
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.2542249469349509, 'optimizer': 'Adam'}
Accuracy:
0.5856784619754255
-----
Total Training Time is (s):
26.36576533317566
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.1339069129108087, 'optimizer': 'Adam'}
Accuracy:
0.5881861242727504
-----
Total Training Time is (s):
25.207025289535522
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.18646617942102145, 'optimizer': 'sgd'}
Accuracy:
0.3290610197868212
-----
Total Training Time is (s):
25.25748372077942
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.2422444255685212, 'optimizer': 'sgd'}
Accuracy:
0.33630537754664286
-----
Total Training Time is (s):
25.461172342300415
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.2448541881379686, 'optimizer': 'sgd'}
Accuracy:
0.31930899972552274
-----
Total Training Time is (s):
26.484083890914917
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.18503897577787976, 'optimizer': 'Adam'}
Accuracy:
0.5792699916494295
-----
Total Training Time is (s):
26.643880128860474
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.32063192160086557, 'optimizer': 'Adam'}
Accuracy:
0.590136528289162
-----
Total Training Time is (s):
27.22426700592041
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.13249574191657534, 'optimizer': 'Adam'}
Accuracy:
0.5859570911200341
-----
Total Training Time is (s):
26.097631216049194
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.29590396234361205, 'optimizer': 'sgd'}
Accuracy:
```

0.33769852327383737

-----  
Total Training Time is (s):

26.8969247341156

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.21040384851779753, 'optimizer': 'Adam'}

Accuracy:

0.5951518529170269

-----  
Total Training Time is (s):

27.601470947265625

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.11880611138425023, 'optimizer': 'Adam'}

Accuracy:

0.5820562830913628

-----  
Total Training Time is (s):

23.013309955596924

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.14665198365052354, 'optimizer': 'Adam'}

Accuracy:

0.46308163833937027

-----  
Total Training Time is (s):

26.464831113815308

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.18015588988489312, 'optimizer': 'sgd'}

Accuracy:

0.3134577876928956

-----  
Total Training Time is (s):

26.698853969573975

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.2799730869754218, 'optimizer': 'sgd'}

Accuracy:

0.309835608804681

-----  
Total Training Time is (s):

23.167685508728027

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.160008848531309, 'optimizer': 'Adam'}

Accuracy:

0.4611312343312625

-----  
Total Training Time is (s):

26.600367546081543

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.16467814646642104, 'optimizer': 'sgd'}

Accuracy:

0.30342713848283676

-----  
Total Training Time is (s):

26.676040172576904

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.34563440284244595, 'optimizer': 'sgd'}

Accuracy:

0.30203399275979415

-----  
Total Training Time is (s):

27.901010513305664

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.26596686122277424, 'optimizer': 'Adam'}

```
{ 'conv_kernel_size': 5, 'dropout_prob': 0.20990001227121, 'optimizer': 'Adam' }
Accuracy:
0.5890220117356393
-----
Total Training Time is (s):
27.972566843032837
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.1641466973375958, 'optimizer': 'Adam'}
Accuracy:
0.5817776539758176
-----
Total Training Time is (s):
28.094647645950317
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.3380391316682997, 'optimizer': 'Adam'}
Accuracy:
0.573697408753107
-----
Total Training Time is (s):
23.941020727157593
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.2156736105684234, 'optimizer': 'Adam'}
Accuracy:
0.45082195598074704
-----
Total Training Time is (s):
24.138392686843872
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.30311341285267246, 'optimizer': 'Adam'}
Accuracy:
0.4522151017037897
-----
Total Training Time is (s):
28.54242205619812
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.31859347865493814, 'optimizer': 'Adam'}
Accuracy:
0.584285316248231
-----
Total Training Time is (s):
28.73721957206726
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.2100158958151668, 'optimizer': 'Adam'}
Accuracy:
0.590136528289162
-----
Total Training Time is (s):
28.72278118133545
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.2108128237693697, 'optimizer': 'Adam'}
Accuracy:
0.5759264419390385
-----
100%|██████████| 25/25 [11:48<00:00, 28.32s/it, best loss: -0.5951518529170269]
=====
Best Hyperparameters {'conv_kernel_size': 2, 'dropout_prob': 0.21040384851779753, 'op
3589/3589 [=====] - 0s 92us/step
=====
Test Accuracy: 0.6032320980816108
```

```

1 # ZCA Whitening
2 from keras.preprocessing.image import ImageDataGenerator
3 from matplotlib import pyplot
4 datagen = ImageDataGenerator()
5
6 #feature standardization
7 x_train = (train_image_array)
8 x_valid = (valid_image_array)
9 x_test = (test_image_array)
10
11 x_train = x_train.reshape((x_train.shape[0], 48, 48, 1))
12 x_valid = x_valid.reshape((x_valid.shape[0], 48, 48, 1))
13 x_test = x_test.reshape((x_valid.shape[0], 48, 48, 1))
14
15 x_train = x_train.astype('float32')
16 x_valid = x_valid.astype('float32')
17 x_test = x_test.astype('float32')
18
19 datagen = ImageDataGenerator(zca_whitening=True)
20 datagen.fit(x_train)
21 datagen.fit(x_valid)
22 datagen.fit(x_test)


1 def optimize_cnn(hyperparameter):
2
3     # Define model using hyperparameters
4     cnn_model = Sequential([Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], ac
5                               Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], activation='relu',
6                               MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
7                               Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu',
8                               Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu',
9                               MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
10                               Flatten(),
11                               Dense(512, activation='relu'),
12                               Dense(7, activation='softmax'),])
13
14     cnn_model.compile(optimizer=hyperparameter['optimizer'], loss='categorical_crossentropy',
15                       time_callback = TimeHistory())
16     hist = cnn_model.fit(x_train, to_categorical(train_labels), epochs=10, batch_size=256)
17     # print(hist.history.accuracy)
18     print("Total Training Time is (s): ", sum(time_callback.times))
19     # Evaluate accuracy on validation data
20     performance = cnn_model.evaluate(x_valid, to_categorical(valid_labels), verbose=0)
21
22     print("Hyperparameters: ", hyperparameter, "Accuracy: ", performance[1])
23     print("-----")
24     # We want to minimize loss i.e. negative of accuracy
25     return({"status": STATUS_OK, "loss": -1*performance[1], "model":cnn_model})
26
27

```



```

27
28 # Define search space for hyper-parameters
29 space = {
30     # The kernel_size for convolutions:
31     'conv_kernel_size': hp.choice('conv_kernel_size', [1, 3, 5]),
32     # Uniform distribution in finding appropriate dropout values
33     'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.35),
34     # Choice of optimizer
35     'optimizer': hp.choice('optimizer', ['Adam', 'sgd']),
36 }
37
38 trials = Trials()
39
40 # Find the best hyperparameters
41 best = fmin(
42     optimize_cnn,
43     space,
44     algo=tpe.suggest,
45     trials=trials,
46     max_evals=25,
47 )
48
49 print("=====")
50 print("Best Hyperparameters", best)
51
52 # You can retrain the final model with optimal hyperparameters on train+validation data
53
54 # Or you can use the model returned directly
55 # Find trial which has minimum loss value and use that model to perform evaluation on test
56 test_model = trials.results[np.argmin([r['loss'] for r in trials.results])]['model']
57
58 performance = test_model.evaluate(x_test, to_categorical(test_labels))
59
60 print("=====")
61 print("Test Accuracy: ", performance[1])

```



```
Total Training Time is (s):
28.13971781730652
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.24814050988778366, 'optimizer': 'sgd'}
Accuracy:
0.34912231820278694
-----
Total Training Time is (s):
29.38590717315674
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.25791132899874547, 'optimizer': 'Adam'}
Accuracy:
0.574811925335693
-----
Total Training Time is (s):
25.396419286727905
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.22730878575841162, 'optimizer': 'Adam'}
Accuracy:
0.4413485650640572
-----
Total Training Time is (s):
23.52994656562805
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.3356007393482103, 'optimizer': 'sgd'}
Accuracy:
0.3170799665977178
-----
Total Training Time is (s):
28.708142280578613
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.11560059994968835, 'optimizer': 'sgd'}
Accuracy:
0.3455001393187242
-----
Total Training Time is (s):
25.343907117843628
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.1569881114299798, 'optimizer': 'Adam'}
Accuracy:
0.45555865143909197
-----
Total Training Time is (s):
23.794260263442993
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.14529594832766432, 'optimizer': 'sgd'}
Accuracy:
0.3232098077500421
-----
Total Training Time is (s):
29.731204509735107
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.14865966844802445, 'optimizer': 'Adam'}
Accuracy:
0.5959877403508524
-----
Total Training Time is (s):
29.16386127471924
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.1668994273753261, 'optimizer': 'sgd'}
Accuracy:
```

0.29757592644190584

-----  
Total Training Time is (s):

24.0995192527771

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.34520159231583475, 'optimizer': 'sgd'}

Accuracy:

0.3037057676274453

-----  
Total Training Time is (s):

25.988826990127563

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.11260333076964793, 'optimizer': 'Adam'}

Accuracy:

0.44552800223733696

-----  
Total Training Time is (s):

24.36750626564026

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.10067291427456249, 'optimizer': 'sgd'}

Accuracy:

0.3240456951838677

-----  
Total Training Time is (s):

29.36147117614746

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.3329262177643334, 'optimizer': 'sgd'}

Accuracy:

0.31707996656865456

-----  
Total Training Time is (s):

24.480132341384888

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.2987819312451138, 'optimizer': 'sgd'}

Accuracy:

0.29311786013647323

-----  
Total Training Time is (s):

30.930524110794067

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.1979817749585424, 'optimizer': 'Adam'}

Accuracy:

0.5812203956616891

-----  
Total Training Time is (s):

29.91033911705017

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.1050616079800136, 'optimizer': 'sgd'}

Accuracy:

0.35218723879348074

-----  
Total Training Time is (s):

29.913865327835083

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.1531927927777971, 'optimizer': 'sgd'}

Accuracy:

0.29172471440512676

-----  
Total Training Time is (s):

30.116098642349243

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.16843832400704467, 'optimizer': 'sgd'}

```
( 'conv_kernel_size': 5, 'dropout_prob': 0.15462614464529809, 'optimizer': 'sgd' }
Accuracy:
0.3507940930787419
-----
Total Training Time is (s):
30.94815754890442
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.15462614464529809, 'optimizer': 'Adam'}
Accuracy:
0.5876288660125967
-----
Total Training Time is (s):
30.283236026763916
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.11775894597544814, 'optimizer': 'sgd'}
Accuracy:
0.2942323767149073
-----
Total Training Time is (s):
31.25113534927368
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.1959757796377455, 'optimizer': 'Adam'}
Accuracy:
0.5748119253315411
-----
Total Training Time is (s):
31.228559494018555
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.1930187501789366, 'optimizer': 'Adam'}
Accuracy:
0.5912510448634442
-----
Total Training Time is (s):
31.54096746444702
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.19116528768022692, 'optimizer': 'Adam'}
Accuracy:
0.5834494288185573
-----
Total Training Time is (s):
31.804988145828247
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.22138591652886783, 'optimizer': 'Adam'}
Accuracy:
0.577598216790082
-----
Total Training Time is (s):
31.88996720314026
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.13687584560602467, 'optimizer': 'Adam'}
Accuracy:
0.5851212037111199
-----
100%|██████████| 25/25 [13:03<00:00, 31.35s/it, best loss: -0.5959877403508524]
=====
Best Hyperparameters {'conv_kernel_size': 2, 'dropout_prob': 0.14865966844802445, 'op
3589/3589 [=====] - 0s 112us/step
=====
Test Accuracy: 0.5918083031568131
```

```

1 # Random Rotations
2 from keras.preprocessing.image import ImageDataGenerator
3 from matplotlib import pyplot
4 datagen = ImageDataGenerator()
5
6 #feature standardization
7 x_train = (train_image_array)
8 x_valid = (valid_image_array)
9 x_test = (test_image_array)
10
11 x_train = x_train.reshape((x_train.shape[0], 48, 48, 1))
12 x_valid = x_valid.reshape((x_valid.shape[0], 48, 48, 1))
13 x_test = x_test.reshape((x_test.shape[0], 48, 48, 1))
14
15 x_train = x_train.astype('float32')
16 x_valid = x_valid.astype('float32')
17 x_test = x_test.astype('float32')
18
19 datagen = ImageDataGenerator(rotation_range=90)
20 datagen.fit(x_train)
21 datagen.fit(x_valid)
22 datagen.fit(x_test)


1 def optimize_cnn(hyperparameter):
2
3     # Define model using hyperparameters
4     cnn_model = Sequential([Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], ac
5                               Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], activation='relu',
6                               MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
7                               Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu',
8                               Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu',
9                               MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
10                               Flatten(),
11                               Dense(512, activation='relu'),
12                               Dense(7, activation='softmax'),])
13
14     cnn_model.compile(optimizer=hyperparameter['optimizer'], loss='categorical_crossentropy',
15                       time_callback = TimeHistory())
16     hist = cnn_model.fit(x_train, to_categorical(train_labels), epochs=10, batch_size=256)
17     # print(hist.history.accuracy)
18     print("Total Training Time is (s): ", sum(time_callback.times))
19     # Evaluate accuracy on validation data
20     performance = cnn_model.evaluate(x_valid, to_categorical(valid_labels), verbose=0)
21
22     print("Hyperparameters: ", hyperparameter, "Accuracy: ", performance[1])
23     print("-----")
24     # We want to minimize loss i.e. negative of accuracy
25     return({"status": STATUS_OK, "loss": -1*performance[1], "model":cnn_model})
26
27

```

```

27
28 # Define search space for hyper-parameters
29 space = {
30     # The kernel_size for convolutions:
31     'conv_kernel_size': hp.choice('conv_kernel_size', [1, 3, 5]),
32     # Uniform distribution in finding appropriate dropout values
33     'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.35),
34     # Choice of optimizer
35     'optimizer': hp.choice('optimizer', ['Adam', 'sgd']),
36 }
37
38 trials = Trials()
39
40 # Find the best hyperparameters
41 best = fmin(
42     optimize_cnn,
43     space,
44     algo=tpe.suggest,
45     trials=trials,
46     max_evals=25,
47 )
48
49 print("=====")
50 print("Best Hyperparameters", best)
51
52 # You can retrain the final model with optimal hyperparameters on train+validation data
53
54 # Or you can use the model returned directly
55 # Find trial which has minimum loss value and use that model to perform evaluation on test
56 test_model = trials.results[np.argmin([r['loss'] for r in trials.results])]['model']
57
58 performance = test_model.evaluate(x_test, to_categorical(test_labels))
59
60 print("=====")
61 print("Test Accuracy: ", performance[1])

```



```
Total Training Time is (s):
31.095835208892822
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.2429661912776022, 'optimizer': 'sgd'}
Accuracy:
0.337977152414294
-----
Total Training Time is (s):
31.687371969223022
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.3226159513161642, 'optimizer': 'sgd'}
Accuracy:
0.3482864307689614
-----
Total Training Time is (s):
28.15788960456848
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.1697117944569278, 'optimizer': 'Adam'}
Accuracy:
0.4446921147993595
-----
Total Training Time is (s):
32.000370264053345
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.21779164306545729, 'optimizer': 'sgd'}
Accuracy:
0.3307327946586242
-----
Total Training Time is (s):
32.85355305671692
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.10663941492259457, 'optimizer': 'sgd'}
Accuracy:
0.35135135135550327
-----
Total Training Time is (s):
32.57329607009888
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.33037373096037836, 'optimizer': 'sgd'}
Accuracy:
0.3379771524433573
-----
Total Training Time is (s):
32.80647110939026
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.1381682853825744, 'optimizer': 'sgd'}
Accuracy:
0.3296182780760382
-----
Total Training Time is (s):
28.958810329437256
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.34275056286314465, 'optimizer': 'Adam'}
Accuracy:
0.4419058233823375
-----
Total Training Time is (s):
32.93373417854309
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.3145589369201526, 'optimizer': 'sgd'}
Accuracy:
```

0.3207021454485654

-----  
Total Training Time is (s):

34.28903603553772

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.33845931234106297, 'optimizer': 'Adam'}

Accuracy:

0.5876288660125967

-----  
Total Training Time is (s):

27.753254175186157

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.2607511117740823, 'optimizer': 'sgd'}

Accuracy:

0.3296182780801901

-----  
Total Training Time is (s):

33.09608292579651

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.13697860925216643, 'optimizer': 'sgd'}

Accuracy:

0.34410699362474484

-----  
Total Training Time is (s):

32.954015254974365

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.10060224254979439, 'optimizer': 'sgd'}

Accuracy:

0.3402061855711622

-----  
Total Training Time is (s):

33.843745470047

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.19430122651964926, 'optimizer': 'Adam'}

Accuracy:

0.5984954026274177

-----  
Total Training Time is (s):

27.96817183494568

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.11444065209952078, 'optimizer': 'sgd'}

Accuracy:

0.3324045695262754

-----  
Total Training Time is (s):

33.4222207069397

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.10390735075834195, 'optimizer': 'sgd'}

Accuracy:

0.3482864307689614

-----  
Total Training Time is (s):

33.40985345840454

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.18057138820927132, 'optimizer': 'sgd'}

Accuracy:

0.32794650321253893

-----  
Total Training Time is (s):

33.81111192703247

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.2277213692352072, 'optimizer': 'Adam'}



```

{'conv_kernel_size': 3, 'dropout_prob': 0.12272133232072, 'optimizer': 'Adam'}
Accuracy:
0.5976595151935922
-----
Total Training Time is (s):
34.45823407173157
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.10651244975702923, 'optimizer': 'Adam'}
Accuracy:
0.5684034550346083
-----
Total Training Time is (s):
34.07155179977417
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.29316343846787873, 'optimizer': 'Adam'}
Accuracy:
0.591529674037116
-----
Total Training Time is (s):
34.26889753341675
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.20573788361981166, 'optimizer': 'Adam'}
Accuracy:
0.5739760379018675
-----
Total Training Time is (s):
34.23374938964844
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.2738452663694213, 'optimizer': 'Adam'}
Accuracy:
0.5982167734828092
-----
Total Training Time is (s):
34.337204694747925
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.27431157338690143, 'optimizer': 'Adam'}
Accuracy:
0.5851212036862085
-----
Total Training Time is (s):
35.21518063545227
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.1917634352204038, 'optimizer': 'Adam'}
Accuracy:
0.5834494288185573
-----
Total Training Time is (s):
35.195539474487305
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.29441982460016747, 'optimizer': 'Adam'}
Accuracy:
0.5639453886959606
-----
100%|██████████| 25/25 [15:16<00:00, 36.66s/it, best loss: -0.5984954026274177]
=====
Best Hyperparameters {'conv_kernel_size': 2, 'dropout_prob': 0.19430122651964926, 'op
3589/3589 [=====] - 0s 123us/step
=====
Test Accuracy: 0.5770409584925612

```

## ▼ FEATURE Design

```
1 from skimage.feature import hog
2 from skimage import data, exposure
3
4 x_train = (train_image_array)
5 x_valid = (valid_image_array)
6 x_test = (test_image_array)
7
8 x_train = x_train.reshape((x_train.shape[0], 48, 48))
9 x_valid = x_valid.reshape((x_valid.shape[0], 48, 48))
10 x_test = x_test.reshape((x_test.shape[0], 48, 48))
11
12 hog_train = []
13 for i in range(len(x_train)):
14     image = x_train[i]
15     fd, hog_image = hog(image, orientations=8, pixels_per_cell=(16, 16),
16                         cells_per_block=(1, 1), visualize=True, multichannel=False)
17     hog_image_rescaled = exposure.rescale_intensity(hog_image, in_range=(0, 10))
18     hog_train.append(hog_image_rescaled)
19
20 hog_valid = []
21 for i in range(len(x_valid)):
22     image = x_valid[i]
23     fd, hog_image = hog(image, orientations=8, pixels_per_cell=(16, 16),
24                         cells_per_block=(1, 1), visualize=True, multichannel=False)
25     hog_image_rescaled = exposure.rescale_intensity(hog_image, in_range=(0, 10))
26     hog_valid.append(hog_image_rescaled)
27
28 hog_test = []
29 for i in range(len(x_test)):
30     image = x_test[i]
31     fd, hog_image = hog(image, orientations=8, pixels_per_cell=(16, 16),
32                         cells_per_block=(1, 1), visualize=True, multichannel=False)
33     hog_image_rescaled = exposure.rescale_intensity(hog_image, in_range=(0, 10))
34     hog_test.append(hog_image_rescaled)
35
36 x_train = np.array(hog_train)
37 x_valid = np.array(hog_valid)
38 x_test = np.array(hog_test)
39
40 x_train = x_train.reshape((x_train.shape[0], 48, 48, 1))
41 x_valid = x_valid.reshape((x_valid.shape[0], 48, 48, 1))
42 x_test = x_test.reshape((x_test.shape[0], 48, 48, 1))
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
```

```

3
4 # Define model using hyperparameters
5 cnn_model = Sequential([Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], ac
6     Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'], activation='relu',
7     MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
8     Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu',
9     Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'], activation='relu',
10    MaxPooling2D(pool_size=(2,2)), Dropout(hyperparameter['dropout_prob']),
11    Flatten(),
12    Dense(512, activation='relu'),
13    Dense(7, activation='softmax'),])
14
15 cnn_model.compile(optimizer=hyperparameter['optimizer'], loss='categorical_crossentropy',
16 time_callback = TimeHistory())
17 hist = cnn_model.fit(x_train, to_categorical(train_labels), epochs=10, batch_size=256)
18 # print(hist.history.accuracy)
19 print("Total Training Time is (s): ", sum(time_callback.times))
20 # Evaluate accuracy on validation data
21 performance = cnn_model.evaluate(x_valid, to_categorical(valid_labels), verbose=0)
22
23 print("Hyperparameters: ", hyperparameter, "Accuracy: ", performance[1])
24 print("-----")
25 # We want to minimize loss i.e. negative of accuracy
26 return({"status": STATUS_OK, "loss": -1*performance[1], "model":cnn_model})
27
28
29 # Define search space for hyper-parameters
30 space = {
31     # The kernel_size for convolutions:
32     'conv_kernel_size': hp.choice('conv_kernel_size', [1, 3, 5]),
33     # Uniform distribution in finding appropriate dropout values
34     'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.35),
35     # Choice of optimizer
36     'optimizer': hp.choice('optimizer', ['Adam', 'sgd']),
37 }
38
39 trials = Trials()
40
41 # Find the best hyperparameters
42 best = fmin(
43     optimize_cnn,
44     space,
45     algo=tpe.suggest,
46     trials=trials,
47     max_evals=25,
48 )
49
50 print("=====")
51 print("Best Hyperparameters", best)
52
53 # You can retrain the final model with optimal hyperparameters on train+validation data
54
55 # Or you can use the model returned directly
56 # Find trial which has minimum loss value and use that model to perform evaluation on test
57 test_model = trials.results[np.argmin([r['loss'] for r in trials.results])]['model']

```

```
58
59 performance = test_model.evaluate(x_test, to_categorical(test_labels))
60
61 print("=====")
62 print("Test Accuracy: ", performance[1])
```



```
Total Training Time is (s):
30.280741214752197
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.15049653462509077, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
26.057782888412476
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.3327359916145542, 'optimizer': 'Adam'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
25.993083477020264
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.2399034427313376, 'optimizer': 'Adam'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
24.862027883529663
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.15711783288164408, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
26.145740032196045
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.13095244695573482, 'optimizer': 'Adam'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
21.854387521743774
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.3277228071731925, 'optimizer': 'Adam'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
21.76009750366211
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.12063233345739746, 'optimizer': 'Adam'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
26.186025619506836
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.1300360838648438, 'optimizer': 'Adam'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
25.091683626174927
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.31162659374425034, 'optimizer': 'sgd'}
Accuracy:
```

0.24937308442878273

-----  
Total Training Time is (s):

26.44800066947937

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.32473622789596435, 'optimizer': 'Adam'}

Accuracy:

0.24937308442878273

-----  
Total Training Time is (s):

25.53569984436035

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.25676837994360124, 'optimizer': 'sgd'}

Accuracy:

0.24937308442878273

-----  
Total Training Time is (s):

25.3596453666687

Hyperparameters:

{'conv\_kernel\_size': 5, 'dropout\_prob': 0.1996249235229065, 'optimizer': 'sgd'}

Accuracy:

0.24937308442878273

-----  
Total Training Time is (s):

26.942335844039917

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.2226626407591457, 'optimizer': 'Adam'}

Accuracy:

0.24937308442878273

-----  
Total Training Time is (s):

25.856187105178833

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.2767054261515769, 'optimizer': 'sgd'}

Accuracy:

0.24937308442878273

-----  
Total Training Time is (s):

25.78175401687622

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.25144256564530276, 'optimizer': 'sgd'}

Accuracy:

0.24937308442878273

-----  
Total Training Time is (s):

21.05579113960266

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.1795332127981628, 'optimizer': 'sgd'}

Accuracy:

0.24937308442878273

-----  
Total Training Time is (s):

21.038676261901855

Hyperparameters:

{'conv\_kernel\_size': 1, 'dropout\_prob': 0.3163431451188834, 'optimizer': 'sgd'}

Accuracy:

0.24937308442878273

-----  
Total Training Time is (s):

26.15276527404785

Hyperparameters:

{'conv\_kernel\_size': 3, 'dropout\_prob': 0.19093288081246018, 'optimizer': 'sgd'}

```

{'conv_kernel_size': 1, 'dropout_prob': 0.21503236289103644, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
21.135244607925415
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.21503236289103644, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
26.280521392822266
Hyperparameters:
{'conv_kernel_size': 3, 'dropout_prob': 0.18627258090208043, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
26.11707830429077
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.15518590423516235, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
21.610968351364136
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.2137888917775868, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
26.30264401435852
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.1607458205645651, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
26.25912308692932
Hyperparameters:
{'conv_kernel_size': 5, 'dropout_prob': 0.10813003008178762, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
Total Training Time is (s):
21.570691347122192
Hyperparameters:
{'conv_kernel_size': 1, 'dropout_prob': 0.27555616085888457, 'optimizer': 'sgd'}
Accuracy:
0.24937308442878273
-----
100%|██████████| 25/25 [10:50<00:00, 26.00s/it, best loss: -0.24937308442878273]
=====
Best Hyperparameters {'conv_kernel_size': 2, 'dropout_prob': 0.15049653462509077, 'op
3589/3589 [=====] - 0s 87us/step
=====
Test Accuracy: 0.2449150181191982

```

